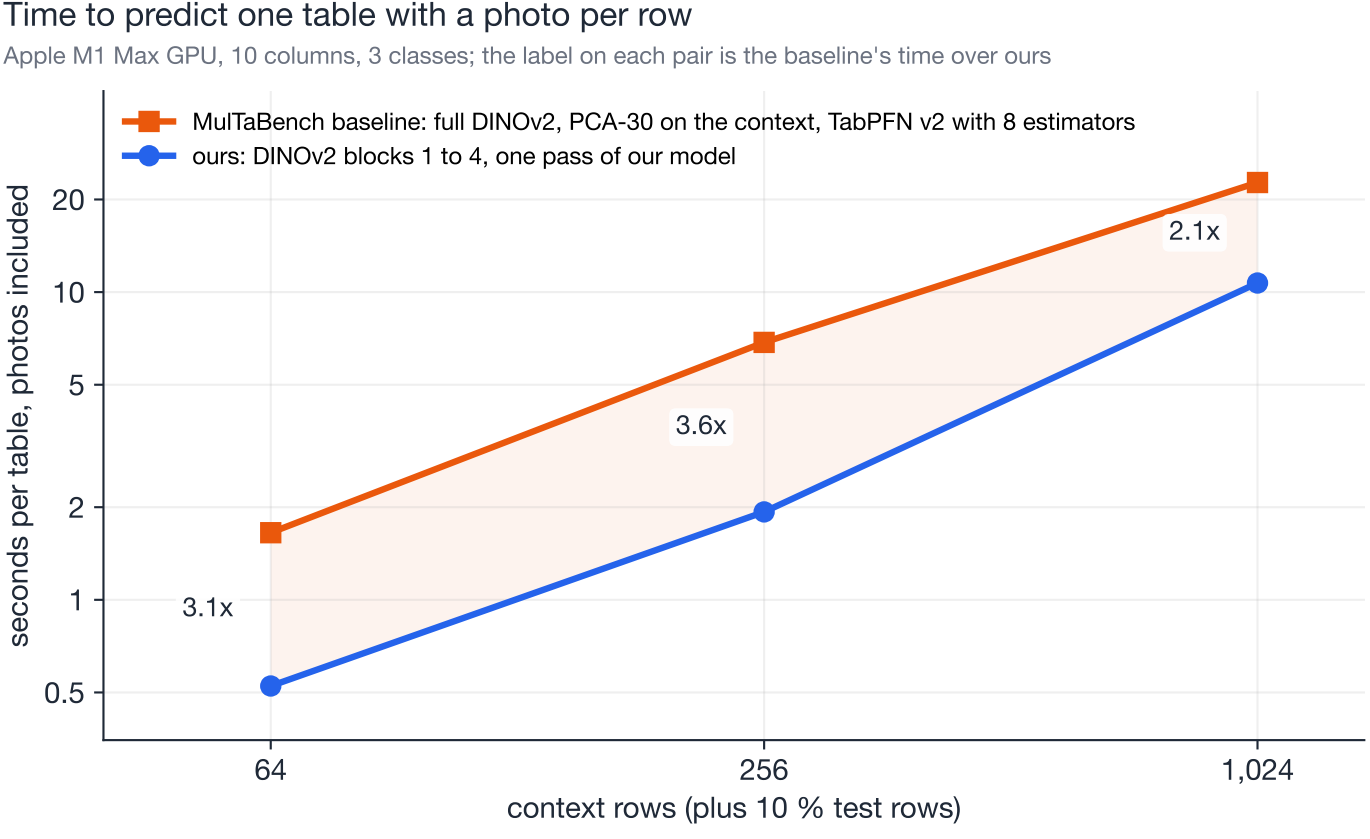
Time to predict one table with a photo per row
Apple M1 Max GPU, 10 columns, 3 classes; the label on each pair is the baseline's time over ours
MulTaBench baseline: full DINOv2, PCA-30 on the context, TabPFN v2 with 8 estimators
ours: DINOv2 blocks 1 to 4, one pass of our model
seconds per table, photos included
20
10
5
2
1
0.5
3.1x
3.6x
2.1x
64
256
1,024
context rows (plus 10 % test rows)